

Disk Allocation

ECE 695

Nov 1

Y. Charlie Hu



1

[lec20] Why Files?



- Physical reality
 - Block oriented
 - Physical sector numbers
 - No protection among users of the system
 - Data might be corrupted if machine crashes
- File system abstraction
 - Byte oriented
 - Named files
 - Users protected from each other
 - Robust to machine failures


"I will save our lab3 solution on platter 5, track 8739, sector 3-4."

"My lab3 solution is in ~ece595/lab3/memory.c" ³

3

Roadmap

- Functionality (API)
 - Basic file system
 - Data structures / disk layout
 - File operations
 - Directories
- Performance
 - Disk allocation
 - Disk scheduling
 - Buffer cache



10

10

Disk Allocation Problem

- Definition: allocate disk blocks when a file is created or grows, and free them when a file is removed or shrinks
- Does this sound familiar?
- How are they different?
 - Granularity?
 - Access performance?
- Shall we approach it like segmentation or paging?

11

11

Disk Allocation Problem



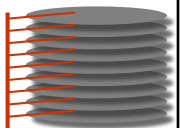
- Two tasks:
 - How to allocate blocks on disk for a file?
 - How to design inode to keep track of blocks?
- Why are they hard?
 1. Disk performance characteristics
 2. Uncertainty about file sizes and access patterns

13

13

Disk Mechanics & Performance



- Platter / Head / Tracks / Sectors / Cylinders
 - Rotation 1000's of RPM (7200, 10k, 15k)
 - Avg seek 5-10 ms
- 
- Assume
 - 255 heads * 38913 tracks * 63 sectors * 512 bytes = 320GB
 - Seek time = 6ms, 7200 RPM → rotational latency = 8ms
 - Block access time = seek time + rotational latency + reading time
 - Accessing a random block: 6ms + 4 ms + 8ms/63 = 638ms/63
 - Accessing the block right after: 8ms/63
 - Implications?

14

14

Disk vs. Memory



Memory

- Latency in 100's of processor cycles
- Transfer rate ~ 1000 MB/s (DDR SDRAM)
- Contiguous allocation gains ~10x
 - Cache hits
 - RAS/CAS (DRAM)

Disk

- Latency in milliseconds
 - 1ms = 10⁶ cycles on 1Ghz machine
- Transfer rate in 30KB/s -- 30MB/s
- Contiguous allocation gains ~100x

15

15

Uncertainty about Files Sizes & Usage Patterns




- File usage determines locality -- How do users access files?
 - Sequential: bytes read in order
 - Random: read/write element out of middle of arrays
 - Whole file or partial file
- How do file sizes vary? (determines inode design)
 - Most files are small
 - Large files use up most of the disk space
 - Large files account for most of the bytes transferred
- Bad news
 - Want FS to be efficient for all cases

16

16

Sequential File Access



17


Demand Paging vs. Disk I/O

18

18

[lec15] Demand Paging algorithms

- Optimal
- FIFO
- FIFO with 2nd chance
- Clock: a simple FIFO with 2nd chance
- Enhanced FIFO with 2nd chance
- NFU
- Approximate LRU




- Definition: pick victim page to swap to disk upon page fault
- Focus on reducing **number** of page faults (going to disk)
 - Don't care which page fault, don't care where on disk
- What kind of *locality* matters?

19

19

Disk Allocation Problem

- Definition: allocate disk blocks when a file is created or grows, and free them when a file is removed or shrinks



- Cannot do anything about number of I/O -- Focus on performance of I/Os (that go to disk), i.e., **where on disk**
 - What kind of locality matters?

20

20

Design goals and expectation



- Optimize I/O performance
 - What can we do for random access patterns?
 - What can we do for sequential access patterns?
- Also want to minimize file header size
 - **File header**: for keeping track of blocks
 - Ideally fit in inode

21

21

Disk Allocation Methods



- Contiguous
- Single-level indexed (Cray-1)
- Linked
- FAT (MS-DOS, OS/2)
- Multi-level indexed (UNIX)

22

22

Contiguous Allocation




- Request in advance for the size of the file
- Search bit map or linked list of free blocks to locate a space
- File header contains
 - first sector number
 - number of sectors

23

23

Contiguous Allocation of Disk Space



directory		
file	start	length
count	0	2
tr	14	3
mail	19	6
list	28	4
f	6	2

Analogy in memory management?

24

Contiguous Allocation



- Request in advance for the size of the file
- Search bit map or linked list to locate a space
- File header contains
 - first sector number
 - number of sectors
- Pros
 - Fast sequential access
 - Easy random access
- Cons
 - External fragmentation
 - Hard to grow files

25

25

Extent-Based Systems




- Later file systems (i.e. Veritas File System – 1st commercial journal FS) use a modified contiguous allocation scheme
- Extent-based file systems allocate disk blocks in **extents**
- An **extent** is a contiguous block of disks
 - Extents are allocated for file allocation
 - A file consists of one or more extents

26

Example: DEMOS (OS for Cray-1, 1977)

- Approach
 - Using contiguous allocation
 - Allow non-contiguous
- Approach
 - 10 (base,size) pointers
 - Indirect for big files
- Pros & cons
 - Can grow (max 10GB)
 - fragmentation
 - Difficult to grow each segment



Analogy in memory management?


27

27

Linked Files




- Approach
 - File header points to 1st block on disk
 - Each block points to next
- Pros
 - Can grow files dynamically
 - No external fragmentation
 - Sequential access easy
- Cons
 - random access: horrible
 - unreliable: losing a block means losing the rest



28

28

Linked Allocation




29

A variation of linked files: File Allocation Table (FAT)

(MS-DOS, OS2 by Microsoft / IBM)


- Approach
 - A section of disk for each partition is reserved
 - One entry for each block
 - A file is a linked list of blocks
 - A directory entry points to the 1st block of the file
- Pros
 - Pros of linked alloc.
 - Simpler than linked alloc.
- Cons
 - Always go to FAT
 - Random access slow for large files



30

Single-Level Indexed Files

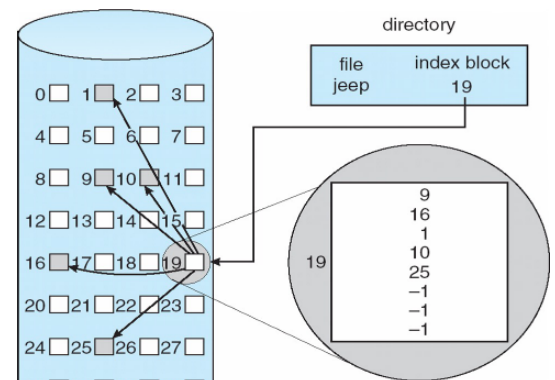
- A user declares max size
- A file header holds an array of pointers to point to disk blocks
- Pros
 - Can grow up to a limit
 - Random access is fast
 - No external fragmentation
- Cons
 - Clumsy to grow beyond limit
 - Up-front declaration a real pain
 - File header cannot fit in inode
 - large inodes would be wasteful



31

31

Example of Indexed Allocation

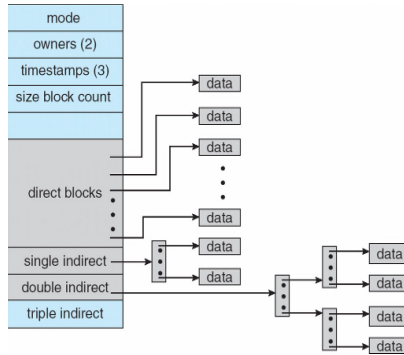


Analogy in memory relocation?

32

Multi-Level Indexed Files (UNIX)

- 13 Pointers in a header
 - 10 direct pointers
 - 11: 1-level indirect
 - 12: 2-level indirect
 - 13: 3-level indirect
- Pros & Cons
 - In favor of small files
 - Can grow
 - Limit is 16G (w/ 1K block)
 - Random access has up to 4 seeks
- How many disk accesses to load block 23, 5, 340?



Analogy in memory management?

Linux ext2

- From Linux kernel documentation for ext2:

"There are pointers to the first 12 blocks which contain the file's data in the inode. There is a pointer to an indirect block (which contains pointers to the next set of blocks), a pointer to a doubly indirect block and a pointer to a trebly indirect block."
- ext2 has 15 pointers.
 - Pointers 1 to 12 point to direct blocks
 - pointer 13 points to an indirect block
 - pointer 14 points to a doubly indirect block
 - pointer 15 points to a trebly indirect block

34

Deep thinking

- Is multi-level indexing a good idea for page table design?

36

Summary

- Seeks kill performance → exploit spatial locality
- Extent-based allocation optimizes sequential access
- Single-level indexed allocation has speed
- Unix file system has great flexibility
- Bitmaps show contiguous free space
- Linked lists easy to search for free blocks

37



Reading

- Chapters 10-11